

## 5. Worksheet: Alpha Diversity

Kelsey Coates; Z620: Quantitative Biodiversity, Indiana University

27 January, 2026

### OVERVIEW

In this exercise, we will explore aspects of local or site-specific diversity, also known as alpha ( $\alpha$ ) diversity. First we will quantify two of the fundamental components of ( $\alpha$ ) diversity: **richness** and **evenness**. From there, we will then discuss ways to integrate richness and evenness, which will include univariate metrics of diversity along with an investigation of the **species abundance distribution (SAD)**.

### Directions:

1. In the Markdown version of this document in your cloned repo, change “Student Name” on line 3 (above) to your name.
2. Complete as much of the worksheet as possible during class.
3. Use the handout as a guide; it contains a more complete description of data sets along with the proper scripting needed to carry out the exercise.
4. Answer questions in the worksheet. Space for your answer is provided in this document and indicated by the “>” character. If you need a second paragraph be sure to start the first line with “>”. You should notice that the answer is highlighted in green by RStudio (color may vary if you changed the editor theme).
5. Before you leave the classroom, **push** this file to your GitHub repo.
6. For the assignment portion of the worksheet, follow the directions at the bottom of this file.
7. When you are done, **Knit** the text and code into a PDF file.
8. After Knitting, submit the completed exercise by creating a **pull request** via GitHub. Your pull request should include this file `AlphaDiversity_Worskheet.Rmd` and the PDF output of `Knitr` (`AlphaDiversity_Worskheet.pdf`).

### 1) R SETUP

In the R code chunk below, please provide the code to: 1) Clear your R environment, 2) Print your current working directory, 3) Set your working directory to your `Week2-Alpha/` folder folder, and 4) Load the `vegan` R package (be sure to install first if you have not already).

```
rm(list = ls())
getwd()
```

```
## [1] "C:/Users/coatesk23/OneDrive - East Carolina University/Documents/QuantBio/QB2026_Coates/weekfour"

setwd("C:/Users/coatesk23/OneDrive - East Carolina University/Documents/QuantBio/QB2026_Coates/weekfour")
```

## 2) LOADING DATA

In the R code chunk below, do the following: 1) Load the BCI dataset, and 2) Display the structure of the dataset (if the structure is long, use the `max.level = 0` argument to show the basic information).

```
library(vegan)

## Warning: package 'vegan' was built under R version 4.5.2

## Loading required package: permute

data(BCI)
str(BCI)

## 'data.frame': 50 obs. of 225 variables:
## $ Abarema.macradenia      : int 0 0 0 0 0 0 0 0 0 1 ...
## $ Vachellia.melanoceras   : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Acalypha.diversifolia    : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Acalypha.macrostachya   : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Adelia.triloba          : int 0 0 0 3 1 0 0 0 5 0 ...
## $ Aegiphila.panamensis    : int 0 0 0 0 1 0 1 0 0 1 ...
## $ Alchornea.costaricensis : int 2 1 2 18 3 2 0 2 2 2 ...
## $ Alchornea.latifolia     : int 0 0 0 0 0 1 0 0 0 0 ...
## $ Alibertia.edulis        : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Allophylus.psilospermus : int 0 0 0 0 1 0 0 0 0 0 ...
## $ Alseis.blackiana         : int 25 26 18 23 16 14 18 14 16 14 ...
## $ Amaioua.corymbosa       : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Anacardium.excelsum     : int 0 0 0 0 0 0 0 1 0 0 ...
## $ Andira.inermis          : int 0 0 0 0 1 1 0 0 1 0 ...
## $ Annona.spraguei         : int 1 0 1 0 0 0 0 1 1 0 ...
## $ Apeiba.glabra           : int 13 12 6 3 4 10 5 4 5 5 ...
## $ Apeiba.tibourbou        : int 2 0 1 1 0 0 0 1 0 0 ...
## $ Aspidosperma.desmanthum : int 0 0 0 1 1 1 0 0 0 1 ...
## $ Astrocaryum.standleyanum: int 0 2 1 5 6 2 2 0 2 1 ...
## $ Astronium.graveolens    : int 6 0 1 3 0 1 2 2 0 0 ...
## $ Attalea.butyracea       : int 0 1 0 0 0 1 1 0 0 0 ...
## $ Banara.guianensis       : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Beilschmiedia.pendula   : int 4 5 7 5 8 6 5 9 11 14 ...
## $ Brosimum.alicastrum    : int 5 2 4 3 2 2 6 4 3 6 ...
## $ Brosimum.guianense     : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Calophyllum.longifolium : int 0 2 0 2 1 2 2 2 2 0 ...
## $ Casearia.aculeata       : int 0 0 0 0 0 0 0 1 0 0 ...
## $ Casearia.arborea        : int 1 1 3 2 4 1 2 3 9 7 ...
## $ Casearia.commersoniana  : int 0 0 1 0 1 0 0 0 1 0 ...
## $ Casearia.guianensis     : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Casearia.sylvestris     : int 2 1 0 0 0 3 1 0 1 1 ...
## $ Cassipourea.guianensis  : int 2 0 1 1 3 4 4 0 2 1 ...
## $ Cavanillesia.platanifolia: int 0 0 0 0 0 0 0 0 0 0 ...
## $ Cecropia.insignis       : int 12 5 7 17 21 4 0 7 2 16 ...
## $ Cecropia.obtusifolia    : int 0 0 0 0 1 0 0 2 0 2 ...
## $ Cedrela.odorata         : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Ceiba.pentandra         : int 0 1 1 0 1 0 0 1 0 1 ...
## $ Celtis.schippii         : int 0 0 0 2 2 0 1 0 0 0 ...
```

```

## $ Cespedesia.spathulata      : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Chamguava.schippiei     : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Chimarrhis.parviflora    : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Maclura.tinctoria       : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Chrysochlamys.eclipes    : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Chrysophyllum.argenteum   : int 4 1 2 2 6 2 3 2 4 2 ...
## $ Chrysophyllum.cainito     : int 0 0 0 0 0 0 1 0 0 0 ...
## $ Coccoloba.coronata       : int 0 0 0 1 2 0 0 1 2 1 ...
## $ Coccoloba.manzinellensis : int 0 0 0 0 0 0 2 0 0 0 ...
## $ Colubrina.glandulosa     : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Cordia.alliodora        : int 2 3 3 7 1 1 2 0 0 2 ...
## $ Cordia.bicolor           : int 12 14 35 23 13 7 5 10 7 13 ...
## $ Cordia.lasiocalyx        : int 8 6 6 11 7 6 6 3 0 4 ...
## $ Coussarea.curvigemma    : int 0 0 0 1 0 2 1 0 1 1 ...
## $ Croton.billbergianus    : int 2 2 0 11 6 0 0 4 2 0 ...
## $ Cupania.cinerea          : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Cupania.latifolia         : int 0 0 0 1 0 0 0 0 0 0 ...
## $ Cupania.rufescens        : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Cupania.seemannii        : int 2 2 1 0 3 0 1 2 2 0 ...
## $ Dendropanax.arboreus     : int 0 3 6 0 5 2 1 6 1 3 ...
## $ Desmopsis.panamensis     : int 0 0 4 0 0 0 0 0 0 1 ...
## $ Diospyros.artanthifolia  : int 1 1 1 1 0 0 0 0 0 1 ...
## $ Dipteryx.oleifera         : int 1 1 3 0 0 0 0 2 1 2 ...
## $ Drypetes.standleyi       : int 2 1 2 0 0 0 0 0 0 0 ...
## $ Elaeis.oleifera          : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Enterolobium.schomburgkii : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Erythrina.costaricensis : int 0 0 0 0 0 3 0 0 1 0 ...
## $ Erythroxylum.macrophyllum: int 0 1 0 0 0 0 0 1 1 1 ...
## $ Eugenia.florida          : int 0 1 0 7 2 0 0 1 1 3 ...
## $ Eugenia.galalonensis     : int 0 0 0 0 0 0 0 1 0 0 ...
## $ Eugenia.nesiotica        : int 0 0 1 0 0 0 5 4 3 0 ...
## $ Eugenia.oerstediana      : int 3 2 5 1 5 2 2 3 3 3 ...
## $ Faramea.occidentalis     : int 14 36 39 39 22 16 38 41 33 42 ...
## $ Ficus.colubrinae         : int 0 1 0 0 0 0 0 0 0 0 ...
## $ Ficus.costaricana        : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Ficus.insipida           : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Ficus.maxima              : int 1 0 0 0 0 0 0 0 0 0 ...
## $ Ficus.obtusifolia        : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Ficus.popenoei            : int 0 0 0 0 0 0 1 0 0 0 ...
## $ Ficus.tonduzii            : int 0 0 1 2 1 0 0 0 0 0 ...
## $ Ficus.trigonata          : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Ficus.yoponensis          : int 1 0 0 0 0 1 1 0 0 0 ...
## $ Garcinia.intermedia       : int 0 1 1 3 2 1 2 2 1 0 ...
## $ Garcinia.madruno          : int 4 0 0 0 1 0 0 0 0 1 ...
## $ Genipa.americana         : int 0 0 1 0 0 0 1 0 1 1 ...
## $ Guapira.myrtiflora        : int 3 1 0 1 1 7 3 1 1 1 ...
## $ Guarea.fuzzy              : int 1 1 0 1 3 0 0 2 0 3 ...
## $ Guarea.grandifolia        : int 0 0 0 0 0 0 0 1 0 0 ...
## $ Guarea.guidonia           : int 2 6 2 5 3 4 4 0 1 5 ...
## $ Guatteria.dumetorum       : int 6 16 6 3 9 7 8 6 2 2 ...
## $ Guazuma.ulmifolia         : int 0 0 0 1 0 0 0 0 0 0 ...
## $ Guettarda.foliacea        : int 1 5 1 2 1 0 0 4 1 3 ...
## $ Gustavia.superba          : int 10 5 0 1 3 1 8 4 4 4 ...
## $ Hampea.appendiculata      : int 0 0 1 0 0 0 0 0 2 1 ...

```

```

## $ Hasseltia.floribunda      : int 5 9 4 11 9 2 7 6 3 4 ...
## $ Heisteria.acuminata     : int 0 0 0 0 1 1 0 0 0 0 ...
## $ Heisteria.concinna       : int 4 5 4 6 4 8 2 5 1 5 ...
## $ Hirtella.americana      : int 0 0 0 0 0 0 0 0 0 0 ...
## $ Hirtella.triandra        : int 21 14 5 4 6 6 7 14 8 7 ...
## $ Hura.crepitans           : int 0 0 0 0 0 2 1 1 0 0 ...
## $ Hieronyma.alchorneoides   : int 0 2 0 0 0 0 0 0 1 0 ...
## [list output truncated]
## - attr(*, "original.names")= chr [1:225] "Abarema.macradenium" "Acacia.melanoceras" "Acalypha.diver...
```

### 3) SPECIES RICHNESS

**Species richness (S)** refers to the number of species in a system or the number of species observed in a sample.

#### Observed richness

In the R code chunk below, do the following:

1. Write a function called `S.obs` to calculate observed richness
2. Use your function to determine the number of species in `site1` of the BCI data set, and
3. Compare the output of your function to the output of the `specnumber()` function in `vegan`.

```

# S.obs <- function(edc) {
#   rowSums(      ) *
# }
S.obs <- function(x = "") {
  rowSums(x > 0) * 1
}

S.obs(BCI[1,])
```

```

## 1
## 93
```

```
specnumber(BCI)
```

```

## 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
## 93 84 90 94 101 85 82 88 90 94 87 84 93 98 93 93 93 89 109 100
## 21 22 23 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40
## 99 91 99 95 105 91 99 85 86 97 77 88 86 92 83 92 88 82 84 80
## 41 42 43 44 45 46 47 48 49 50
## 102 87 86 81 81 86 102 91 91 93
```

```
specnumber(BCI[1,])
```

```

## 1
## 93
```

```
specnumber(BCI[1:4,])
```

```
## 1 2 3 4  
## 93 84 90 94
```

**Question 1:** Does `specnumber()` from `vegan` return the same value for observed richness in `site1` as our function `S.obs`? What is the species richness of the first four sites (i.e., rows) of the BCI matrix? > 93 84 90 94 >

**Answer 1:** >The species number function in vegan returned the same value.

### Coverage: How well did you sample your site?

In the R code chunk below, do the following:

1. Write a function to calculate Good's Coverage, and
2. Use that function to calculate coverage for all sites in the BCI matrix.

```
C <- function(x = ""){  
  1 - (rowSums(x == 1) / rowSums(x))  
}  
  
C(BCI)  
  
##          1         2         3         4         5         6         7         8  
## 0.9308036 0.9287356 0.9200864 0.9468504 0.9287129 0.9174757 0.9326923 0.9443155  
##          9        10        11        12        13        14        15        16  
## 0.9095355 0.9275362 0.9152120 0.9071038 0.9242054 0.9132420 0.9350649 0.9267735  
##         17        18        19        20        21        22        23        24  
## 0.8950131 0.9193084 0.8891455 0.9114219 0.8946078 0.9066986 0.8705882 0.9030612  
##         25        26        27        28        29        30        31        32  
## 0.9095023 0.9115479 0.9088729 0.9198966 0.8983516 0.9221053 0.9382423 0.9411765  
##         33        34        35        36        37        38        39        40  
## 0.9220183 0.9239374 0.9267887 0.9186047 0.9379310 0.9306488 0.9268868 0.9386503  
##         41        42        43        44        45        46        47        48  
## 0.8880597 0.9299517 0.9140049 0.9168704 0.9234234 0.9348837 0.8847059 0.9228916  
##         49        50  
## 0.9086651 0.9143519
```

**Question 2:** Answer the following questions about coverage:

- a. What is the range of values that can be generated by Good's Coverage?
- b. What would we conclude from Good's Coverage if  $n_i$  equaled  $N$ ?
- c. What portion of taxa in `site1` was represented by singletons?
- d. Make some observations about coverage at the BCI plots.

**Answer 2a:**  $0 < x < 1$

**Answer 2b:** That would mean all the species sampled were singletons.

**Answer 2c:** 0.93

**Answer 2d:** With a Good's coverage close to 1 we can conclude the sampling coverage was sufficient to capture most of the biodiversity in the community. Each BCI plot had a coverage > 88 which is good coverage.

## Estimated richness

In the R code chunk below, do the following:

1. Load the microbial dataset (located in the Week2-Alpha/data folder),
2. Transform and transpose the data as needed (see handout),
3. Create a new vector (`soilbac1`) by indexing the bacterial OTU abundances of any site in the dataset,
4. Calculate the observed richness at that particular site, and
5. Calculate coverage of that site

```
soilbac <- read.table("soilbac.txt", sep = "\t", header = TRUE, row.names = 1)
soilbac.t <- as.data.frame(t(soilbac))
soilbac1 <- soilbac.t[1,]

S.obs(soilbac1)

## T1_1
## 1074

C(soilbac1)

##      T1_1
## 0.6479471
```

**Question 3:** Answer the following questions about the soil bacterial dataset.

- a. How many sequences did we recover from the sample `soilbac1`, i.e.  $N$ ?
- b. What is the observed richness of `soilbac1`?
- c. How does coverage compare between the BCI sample (`site1`) and the KBS sample (`soilbac1`)?

**Answer 3a:** 1074

**Answer 3b:** 0.648

**Answer 3c:** Sampling coverage was higher in `site1` of the BCI plots than the KBS site(s).

## Richness estimators

In the R code chunk below, do the following:

1. Write a function to calculate **Chao1**,
2. Write a function to calculate **Chao2**,
3. Write a function to calculate **ACE**, and
4. Use these functions to estimate richness at `site1` and `soilbac1`.

```

S.chao1 <- function(x = ""){
  S.obs(x) + (sum(x == 1)^2) / (2 * sum(x == 2))
}

S.chao2 <- function(site = "", SbyS = ""){
  SbyS = as.data.frame(SbyS)
  x = SbyS[site, ]
  SbyS.pa <- (SbyS > 0) * 1 # convert the SbyS to presence/absence
  Q1 = sum(colSums(SbyS.pa) == 1) # species observed once
  Q2 = sum(colSums(SbyS.pa) == 2) # species observed twice
  S.chao2 = S.obs(x) + (Q1^2)/(2 * Q2)
  return(S.chao2)
}

S.ace <- function(x = "", thresh = 10){
  x <- x[x>0] # excludes zero-abundance taxa
  S.abund <- length(which(x > thresh)) # richness of abundant taxa
  S.rare <- length(which(x <= thresh)) # richness of rare taxa
  singlt <- length(which(x == 1)) # number of singleton taxa
  N.rare <- sum(x[which(x <= thresh)]) # abundance of rare individuals
  C.ace <- 1 - (singlt / N.rare) # coverage (prop non-singlt rare inds)
  i <- c(1:thresh) # threshold abundance range
  count <- function(i, y){ # counter to go through i range
    length(y[y == i])
  }
  a.1 <- sapply(i, count, x) # number of individuals in richness i richness classes
  f.1 <- (i * (i - 1)) * a.1 # k(k-1)kf sensu Gotelli
  G.ace <- (S.rare/C.ace)*(sum(f.1)/(N.rare*(N.rare-1)))
  S.ace <- S.abund + (S.rare/C.ace) + (singlt/C.ace) * max(G.ace,0)
  return(S.ace)
}

S.chao1(BCI[1,])

##          1
## 119.6944

S.chao1(soilbac1)

##          T1_1
## 2628.514

# S.chao2(BCI[1,]) #Error in xj[i] : invalid subscript type 'list'
S.chao2(site = 1, SbyS = BCI) # site by species

##          1
## 104.6053

S.chao2(site=1, SbyS=soilbac.t)

```

```

##      T1_1
## 21055.39

S.ace(BCI[1,])

## [1] 159.3404

S.ace(soilbac1)

## [1] 4465.983

```

**Question 4:** What is the difference between ACE and the Chao estimators? Do the estimators give consistent results? Which one would you choose to use and why?

**Answer 4:** The Chao estimator The estimators did not give consistent results for the different data sets. Chao1 I would use the

### Rarefaction

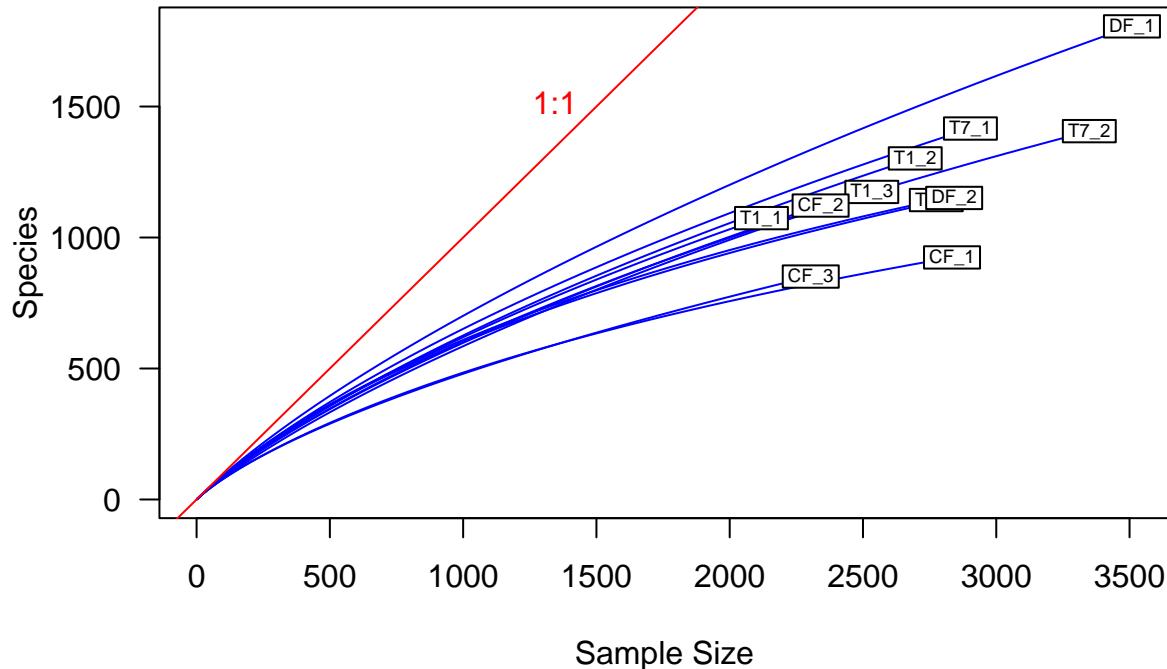
In the R code chunk below, please do the following:

1. Calculate observed richness for all samples in `soilbac`,
2. Determine the size of the smallest sample,
3. Use the `rarefy()` function to rarefy each sample to this level,
4. Plot the rarefaction results, and
5. Add the 1:1 line and label.

```

soilbac.S <- S.obs(soilbac.t)
min.N <- min(rowSums(soilbac.t))
S.rarefy <- rarefy(x = soilbac.t, sample = min.N, se = TRUE)
rarecurve(x = soilbac.t, step = 20, col = "blue", cex = 0.6, las = 1)
abline(0, 1, col = 'red')
text(1500, 1500, "1:1", pos = 2, col = 'red')

```



#### 4) SPECIES EVENNESS

Here, we consider how abundance varies among species, that is, **species evenness**.

##### Visualizing evenness: the rank abundance curve (RAC)

One of the most common ways to visualize evenness is in a **rank-abundance curve** (sometime referred to as a rank-abundance distribution or Whittaker plot). An RAC can be constructed by ranking species from the most abundant to the least abundant without respect to species labels (and hence no worries about ‘ties’ in abundance).

In the R code chunk below, do the following:

1. Write a function to construct a RAC,
2. Be sure your function removes species that have zero abundances,
3. Order the vector (RAC) from greatest (most abundant) to least (least abundant), and
4. Return the ranked vector

```
RAC <- function(x = ""){
  x.ab = x[x > 0]
  x.ab.ranked = x.ab[order(x.ab, decreasing = TRUE)]
  as.data.frame(lapply(x.ab.ranked, unlist))
```

```

    return(x.ab.ranked)
}

```

Now, let us examine the RAC for `site1` of the BCI data set.

In the R code chunk below, do the following:

1. Create a sequence of ranks and plot the RAC with natural-log-transformed abundances,
2. Label the x-axis “Rank in abundance” and the y-axis “log(abundance)”

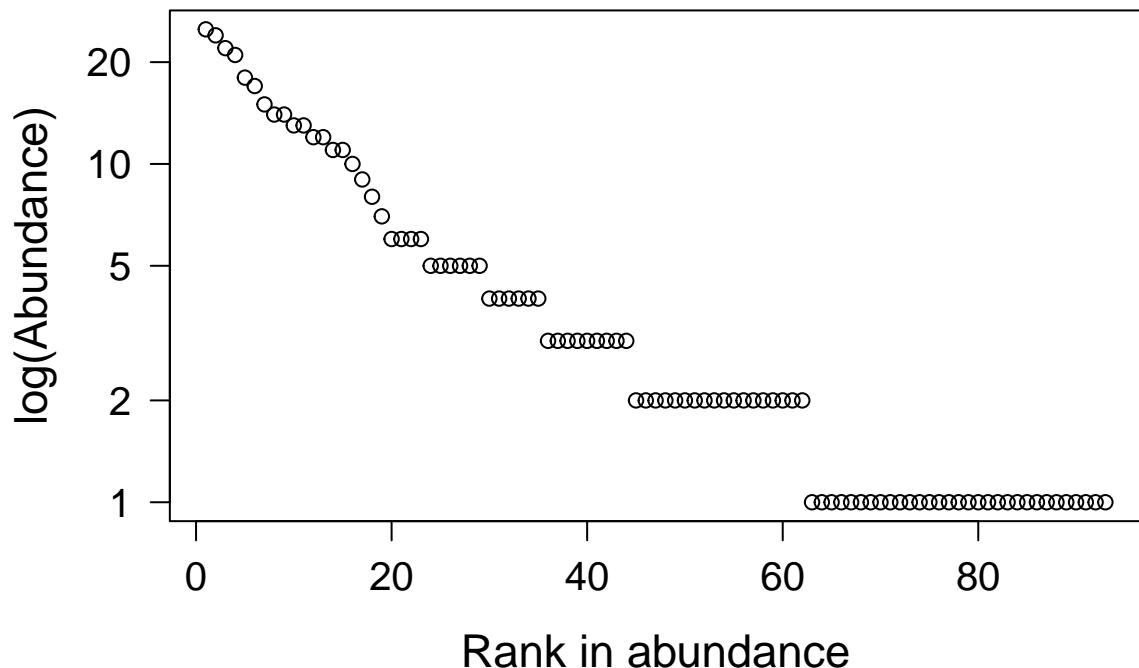
```

plot.new()
site1 <- BCI[1, ]

rac <- RAC(x = site1)
ranks <- as.vector(seq(1, length(rac)))
opar <- par(no.readonly = TRUE)                      # Saves default plot parameters
par(mar = c(5.1, 5.1, 4.1, 2.1))                   # New settings for par
plot(ranks, log(rac), type = 'p', axes = F,          # Plots w/o axes
     xlab = "Rank in abundance", ylab = "log(Abundance)", 
     las = 1, cex.lab = 1.4, cex.axis = 1.25)

box()                                              # Manually adds border
axis(side = 1, labels = T, cex.axis = 1.25)          # Manually adds X-axis
axis(side = 2, las = 1, cex.axis = 1.25,             # Manually adds Log-Scaled Y-axis
     labels = c(1, 2, 5, 10, 20), at = log(c(1, 2, 5, 10, 20)))

```



```
par <- opar
```

**Question 5:** What effect does visualizing species abundance data on a log-scaled axis have on how we interpret evenness in the RAC?

**Answer 5:** Adjusting the scale to log makes species abundance look more even than they are. Without changing to log scales dominant species would be emphasized and rare species would be smushed near the axis.

Now that we have visualized unevenness, it is time to quantify it using Simpson's evenness ( $E_{1/D}$ ) and Smith and Wilson's evenness index ( $E_{var}$ ).

### Simpson's evenness ( $E_{1/D}$ )

In the R code chunk below, do the following:

1. Write the function to calculate  $E_{1/D}$ , and
2. Calculate  $E_{1/D}$  for `site1`.

```
SimpE <- function(x = ""){  
  S <- S.obs(x)  
  x = as.data.frame(x)  
  D <- diversity(x, "inv")  
  E <- (D)/S  
  return(E)  
}  
  
site1 <- BCI[1,]  
SimpE(site1)
```

```
##           1  
## 0.4238232
```

### Smith and Wilson's evenness index ( $E_{var}$ )

In the R code chunk below, please do the following:

1. Write the function to calculate  $E_{var}$ ,
2. Calculate  $E_{var}$  for `site1`, and
3. Compare  $E_{1/D}$  and  $E_{var}$ .

```
Evar <- function(x){  
  x <- as.vector(x[x > 0])  
  1 - (2/pi) * atan(var(log(x)))  
}  
  
Evar(site1)
```

```
## [1] 0.5067211
```

**Question 6:** Compare estimates of evenness for `site1` of BCI using  $E_{1/D}$  and  $E_{var}$ . Do they agree? If so, why? If not, why? What can you infer from the results.

*Answer 6:*

Evar was slightly greater than Simpson's, suggesting they are in agreement. These results suggest there is moderate evenness in the community.

## 5) INTEGRATING RICHNESS AND EVENNESS: DIVERSITY METRICS

So far, we have introduced two primary aspects of diversity, i.e., richness and evenness. Here, we will use popular indices to estimate diversity, which explicitly incorporate richness and evenness. We will write our own diversity functions and compare them against the functions in `vegan`.

**Shannon's diversity (a.k.a., Shannon's entropy)**

In the R code chunk below, please do the following:

1. Provide the code for calculating  $H'$  (Shannon's diversity),
2. Compare this estimate with the output of `vegan`'s diversity function using `method = "shannon"`.

```
ShanH <- function(x = ""){  
  H = 0  
  for (n_i in x){  
    if(n_i > 0) {  
      p = n_i / sum(x)  
      H = H - p*log(p)  
    }  
  }  
  return(H)  
}  
  
ShanH(site1)
```

```
## [1] 4.018412
```

```
diversity(site1, index = "shannon")
```

```
## [1] 4.018412
```

## Simpson's diversity (or dominance)

In the R code chunk below, please do the following:

1. Provide the code for calculating D (Simpson's diversity),
2. Calculate both the inverse ( $1/D$ ) and  $1 - D$ ,
3. Compare this estimate with the output of `vegan`'s `diversity` function using method = "simp".

```
SimpD <- function(x = ""){  
  D = 0  
  N = sum(x)  
  for (n_i in x){  
    D = D + (n_i^2)/(N^2)  
  }  
  return(D)  
}
```

```
D.inv <- 1/SimpD(site1)  
D.sub <- 1-SimpD(site1)
```

```
D.inv
```

```
## [1] 39.41555
```

```
D.sub
```

```
## [1] 0.9746293
```

```
diversity(site1, "inv")
```

```
## [1] 39.41555
```

```
diversity(site1, "simp")
```

```
## [1] 0.9746293
```

## Fisher's $\alpha$

In the R code chunk below, please do the following:

1. Provide the code for calculating Fisher's  $\alpha$ ,
2. Calculate Fisher's  $\alpha$  for `site1` of BCI.

```
rac <- as.vector(site1[site1 > 0])  
invD <- diversity(rac, "inv")  
invD
```

```
## [1] 39.41555
```

```
Fisher <- fisher.alpha(rac)
Fisher
```

```
## [1] 35.67297
```

```
ShanH(site1)
```

```
## [1] 4.018412
```

**Question 7:** How is Fisher's  $\alpha$  different from  $E_{H'}$  and  $E_{var}$ ? What does Fisher's  $\alpha$  take into account that  $E_{H'}$  and  $E_{var}$  do not?

**Answer 7:** Fisher's alpha is not an evennes metric like  $E_{H'}$  and  $E_{var}$ , Fisher's alpha is a diversity metric.

## 6) HILL NUMBERS

Remember that we have learned about the advantages of Hill Numbers to measure and compare diversity among samples. We also learned to explore the effects of rare species in a community by examining diversity for a series of exponents  $q$ .

**Question 8:** Using `site1` of BCI and `vegan` package, a) calculate Hill numbers for  $q$  exponent 0, 1 and 2 (richness, exponential Shannon's entropy, and inverse Simpson's diversity). b) Interpret the effect of rare species in your community based on the response of diversity to increasing exponent  $q$ .

```
# Hill number q = 0: Species richness
D0 <- specnumber(site1)

# Hill number q = 1: Exponential Shannon entropy
H <- diversity(site1, index = "shannon")
D1 <- exp(H)

# Hill number q = 2: Inverse Simpson diversity
lambda <- diversity(site1, index = "simpson")
D2 <- 1 / lambda

D0

## 1
## 93

D1

## [1] 55.6127

D2

## [1] 1.026031
```

**Answer 8a:** 93,55.6, 1.02 **Answer 8b:** The community contains many rare species, but diversity is dominated by a smaller set of abundant species as  $q$  increases. As  $q$  increases, the weight of rare species decreases, meaning the abundance of rare species is not overestimated.

## 7) MOVING BEYOND UNIVARIATE METRICS OF $\alpha$ DIVERSITY

The diversity metrics that we just learned about attempt to integrate richness and evenness into a single, univariate metric. Although useful, information is invariably lost in this process. If we go back to the rank-abundance curve, we can retrieve additional information – and in some cases – make inferences about the processes influencing the structure of an ecological system.

### Species abundance models

The RAC is a simple data structure that is both a vector of abundances. It is also a row in the site-by-species matrix (minus the zeros, i.e., absences).

Predicting the form of the RAC is the first test that any biodiversity theory must pass and there are no less than 20 models that have attempted to explain the uneven form of the RAC across ecological systems.

In the R code chunk below, please do the following:

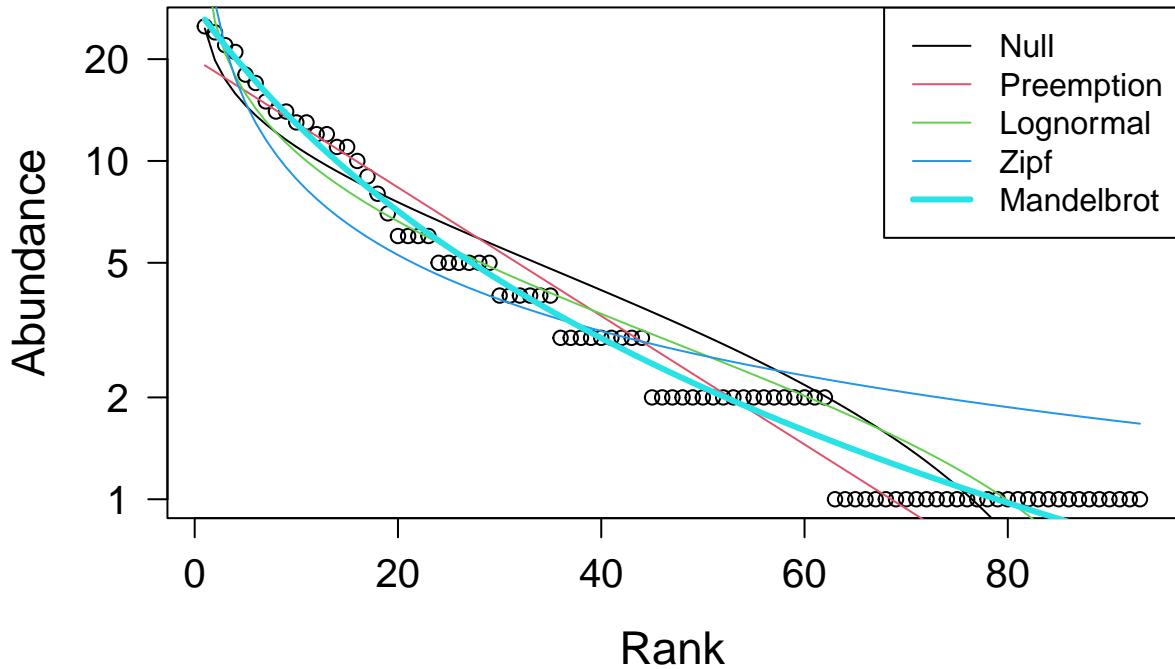
1. Use the `radfit()` function in the `vegan` package to fit the predictions of various species abundance models to the RAC of `site1` in BCI,
2. Display the results of the `radfit()` function, and
3. Plot the results of the `radfit()` function using the code provided in the handout.

```
RACresults <- radfit(site1)
```

```
RACresults
```

```
##  
## RAD models, family poisson  
## No. of species 93, total abundance 448  
##  
##           par1      par2      par3    Deviance   AIC      BIC  
## Null          39.5261 315.4362 315.4362  
## Preemption  0.042797 21.8939 299.8041 302.3367  
## Lognormal    1.0687  1.0186  25.1528 305.0629 310.1281  
## Zipf         0.11033 -0.74705  61.0465 340.9567 346.0219  
## Mandelbrot  100.52   -2.312   24.084  4.2271 286.1372 293.7350
```

```
plot.new()  
plot(RACresults, las = 1, cex.lab = 1.4, cex.axis = 1.25)
```



**Question 9:** Answer the following questions about the rank abundance curves: a) Based on the output of `radfit()` and plotting above, discuss which model best fits our rank-abundance curve for `site1`? b) Can we make any inferences about the forces, processes, and/or mechanisms influencing the structure of our system, e.g., an ecological community?

**Answer 9a:** The Mandelbrot model fit the rank-abundance curve the best with both the lowest BIC and AIC. **Answer 9b:** The Mandelbrot model fitting the best suggests that there may be strong competitive forces in the community.

**Question 10:** Answer the following questions about the preemption model: a. What does the preemption model assume about the relationship between total abundance ( $N$ ) and total resources that can be preempted? b. Why does the niche preemption model look like a straight line in the RAD plot?

**Answer 10a:** The preemption model tells us that species abundance is directly related to resource availability and species will sequentially preempt the fraction of remaining resources until the resources run out. **Answer 10b:** The model looks like a straight line because as available resources decrease, so does the number of species reproducing. The log scale makes the line look straight.

**Question 10:** Why is it important to account for the number of parameters a model uses when judging how well it explains a given set of data?

**Answer 11:** A model with many parameters could be overfit and thus tell us less about the results (meaning the ecological context is lost by fitting the noise), sometimes the best model is the one with the best fit with the least parameters.

## SYNTHESIS

- As stated by Magurran (2004) the  $D = \sum p_i^2$  derivation of Simpson's Diversity only applies to communities of infinite size. For anything but an infinitely large community, Simpson's Diversity index is calculated as  $D = \sum \frac{n_i(n_i-1)}{N(N-1)}$ . Assuming a finite community, calculate Simpson's D, 1 - D, and Simpson's inverse (i.e. 1/D) for site 1 of the BCI site-by-species matrix.

```
SimpD(site1)
```

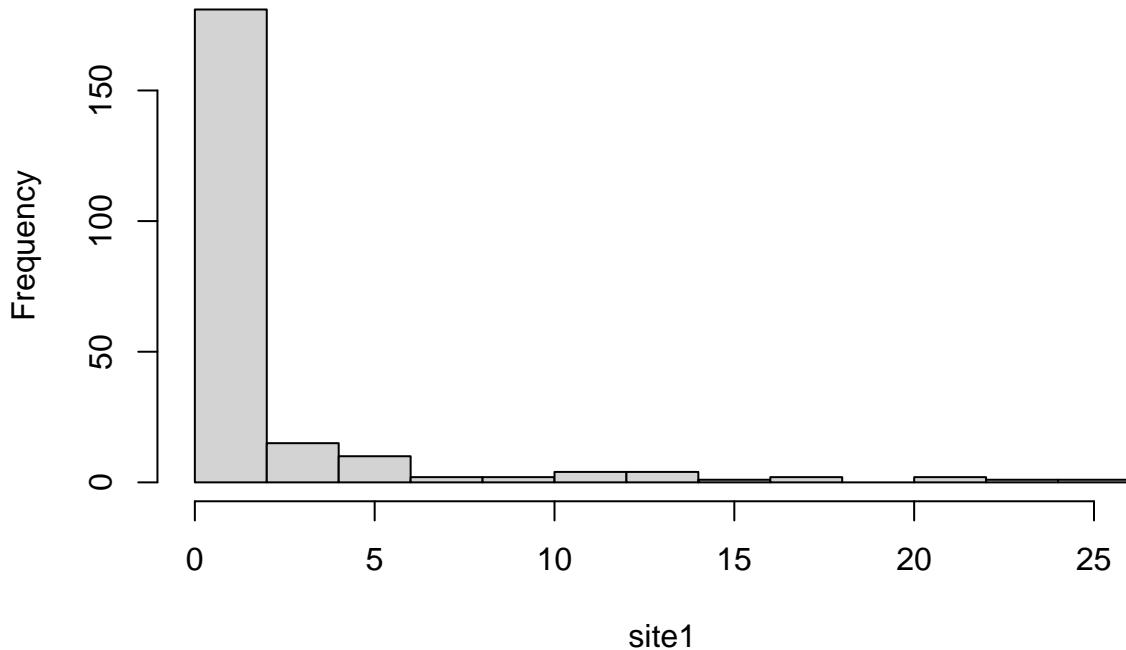
```
## [1] 0.0253707
```

```
D.inv <- 1/SimpD(site1)
D.sub <- 1-SimpD(site1)
```

- Along with the rank-abundance curve (RAC), another way to visualize the distribution of abundance among species is with a histogram (a.k.a., frequency distribution) that shows the frequency of different abundance classes. For example, in a given sample, there may be 10 species represented by a single individual, 8 species with two individuals, 4 species with three individuals, and so on. In fact, the rank-abundance curve and the frequency distribution are the two most common ways to visualize the species-abundance distribution (SAD) and to test species abundance models and biodiversity theories. To address this homework question, use the R function **hist()** to plot the frequency distribution for site 1 of the BCI site-by-species matrix, and describe the general pattern you see.

```
site1 <- as.numeric(site1)
hist(site1)
```

**Histogram of site1**



> This histogram shows that site one of the BCI dataset has many rare species

3. We asked you to find a biodiversity dataset with your partner. This data could be one of your own or it could be something that you obtained from the literature. Load that dataset. How many sites are there? How many species are there in the entire site-by-species matrix? Any other interesting observations based on what you learned this week?

```
data(dune)
nrow(dune) # Number of sites

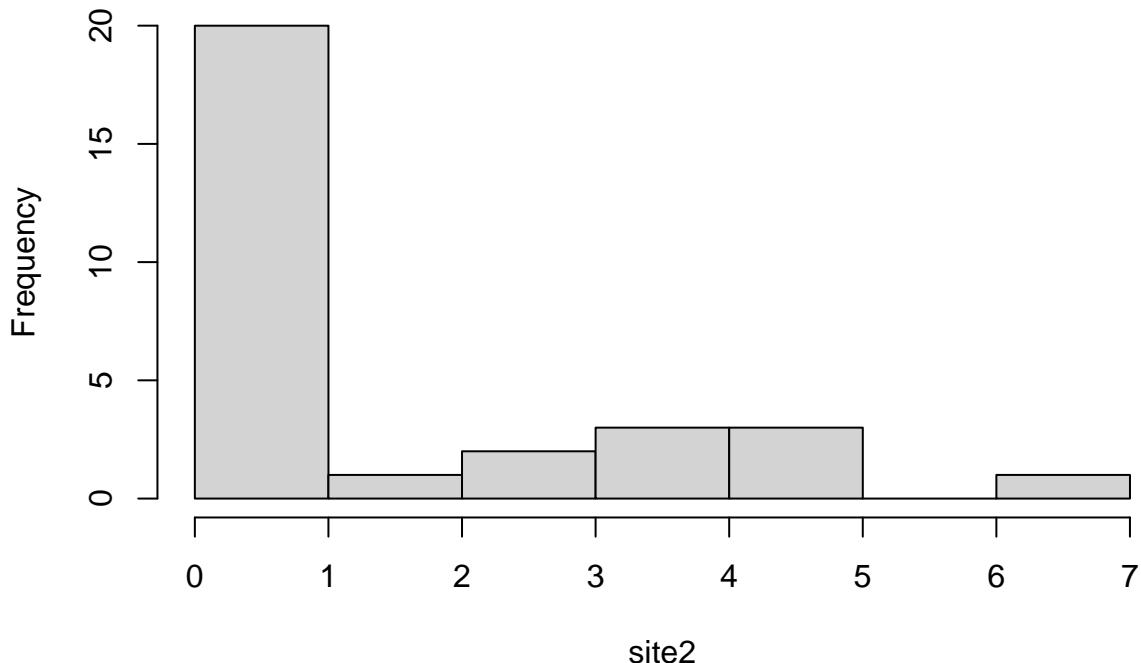
## [1] 20

ncol(dune) # Number of species

## [1] 30

site2 <- as.numeric(dune[2,])
hist(site2)
```

**Histogram of site2**



> The second site of the dune dataset has many rare species.

## SUBMITTING YOUR ASSIGNMENT

Use Knitr to create a PDF of your completed 5.AlphaDiversity\_Worksheet.Rmd document, push it to GitHub, and create a pull request. Please make sure your updated repo include both the pdf and RMarkdown files. Unless otherwise noted, this assignment is due on **Wednesday, January 29<sup>th</sup>, 2025 at 12:00 PM (noon)**.